

Opt performance

This document summarizes the main observations about the performance of the implementation of the Regularized Multinomial Regression in `mtool`. This implementation uses a Stochastic Variance Reduced Gradient (SVRG) algorithm in order to solve the proximal operator.

1 Generative Data Model —

Data are simulated from a $K = 2$ competing risks proportional hazards model. The cause-specific hazard for cause $k \in \{1, 2\}$ for individual i with covariates X_i is $\lambda_k(t|X_i) = \lambda_{0k}(t) \exp(X_i^T \beta_k)$. The baseline hazards $\lambda_{0k}(t)$ follow a Weibull distribution $\lambda_{0k}(t) = h_k \gamma_k t^{\gamma_k - 1}$, with parameters $(h_1, \gamma_1) = (0.55, 1.5)$ and $(h_2, \gamma_2) = (0.05, 1.5)$, implying increasing baseline hazards, higher for cause 1.

The coefficient vectors $\beta_1, \beta_2 \in \mathbb{R}^p$ are sparse. Only the first 10 covariates X_1, \dots, X_{10} have non-zero effects: $\beta_{1j} = +1$ and $\beta_{2j} = -1$ for $j = 1, \dots, 10$, with all other $\beta_{kj} = 0$. These covariates increase $\lambda_1(t|X_i)$ and decrease $\lambda_2(t|X_i)$.

Event times T_i and causes C_i are generated by simulating potential failure times T_{ik} from $\lambda_k(t|X_i)$ and setting $T_i = \min(T_{i1}, T_{i2})$ with C_i being the index yielding the minimum. Independent censoring times $T_{cens,i}$ (rate 0.25) are generated. Observed data consist of $(ftime_i, fstatus_i)$, where $ftime_i = \min(T_i, T_{cens,i})$ and the status $fstatus_i = C_i \cdot \mathbb{1}(T_i \leq T_{cens,i})$ (with $fstatus_i = 0$ indicating censoring).

```
#####  
# ' Create the case-base sampled dataset  
# '  
# ' @param surv_obj  
# ' @param cov_matrix  
# ' @param ratio  
# ' @return List of 4 elements, containing the necessary data to fit a case-base  
# ' regression model.  
create_cbDataset <- function(surv_obj, cov_matrix, ratio = 5) {  
  n <- nrow(surv_obj)  
  B <- sum(surv_obj[, "time"])  
  c1 <- sum(surv_obj[, "status"] == 1)  
  c2 <- sum(surv_obj[, "status"] == 2)  
  b <- ratio * (c1)  
  offset <- log(B/b)  
  probb_select <- surv_obj[, "time"]/B  
  # Create base series
```

```

which_pm <- sample(n, b, replace = TRUE, prob = prob_select)
bSeries <- as.matrix(surv_obj[which_pm, ])
time_bseries <- runif(b) * bSeries[, "time"]
cov_bseries <- cov_matrix[which_pm, , drop = FALSE]
event_bseries <- rep(0L, nrow(bSeries))
# Extract case series
cSeries <- as.matrix(surv_obj[surv_obj[, "status"] != 0L, ])
time_cseries <- cSeries[, "time"]
cov_cseries <- cov_matrix[surv_obj[, "status"] != 0L, , drop = FALSE]
event_cseries <- cSeries[, "status"]
# Combine and return
output <- list("time" = c(time_bseries, time_cseries),
              "event_ind" = c(event_bseries, event_cseries),
              "covariates" = rbind(cov_bseries, cov_cseries),
              "offset" = rep(offset, nrow(bSeries) + nrow(cSeries)))

return(output)
}

##### Function to compute weibull hazard #####
weibull_hazard <- Vectorize(function(gamma, lambda, t) {
  return(gamma * lambda * t^(gamma - 1))
})

##### Function to simulate from cause-specific hazards #####
cause_hazards_sim <- function(p, n, beta1, beta2,
                              nblocks = 4, cor_vals = c(0.7, 0.4, 0.6, 0.5),
                              num.true = 20, h1 = 0.55, h2 = 0.10,
                              gamma1 = 100, gamma2 = 100, max_time = 1.5,
                              noise_cor = 0.1,
                              rate_cens = 0.05, min_time = 0.002,
                              exchangeable = FALSE) {

# Warnings
if(length(beta1) != length(beta2)) stop("Dimension of beta1 and beta2 should be the same")
if(nblocks != length(cor_vals)) stop("Dim of nblocks and corr for blocks should match")
if(isTRUE(exchangeable)) {
  # Create an empty matrix
  mat <- matrix(noise_cor, nrow = p, ncol = p)
  # Set the correlation values
  cor_exchangeable <- 0.5
  # Set the upper triangular and lower triangular parts
  mat[1:num.true, 1:num.true] <- cor_exchangeable
  # Print the matrix
  diag(mat) <- rep(1, length(diag(mat)))
  X <- mvtnorm::rmvnorm(n, mean = rep(0, p), sigma = mat)
} else {
  # Set the number of variables per block
  vpb <- num.true/nblocks

```

```

# Set the correlation values for each covariate block
correlation_values <- cor_vals
# Initialize empty matrix
correlation_matrix <- matrix(noise_cor, nrow = p, ncol = p)
# Generate the covariance matrix with block correlations
for (i in 1:nblocks) {
  start_index <- (i - 1) * vpb + 1
  end_index <- i * vpb
  correlation_matrix[start_index:end_index, start_index:end_index] <- correlation_values[i]
}
# Diagonal elements should be 1
diag(correlation_matrix) <- rep(1, length(diag(correlation_matrix)))
X <- mvtnorm::rmvnorm(n, mean = rep(0, p), sigma = correlation_matrix)
}
X <- as.matrix(X)
# Specify rate parameters
lambda1 <- h1 * exp(X %>% beta1)
lambda2 <- h2 * exp(X %>% beta2)
# Define cdf - U
cdf_U <- function(t, gamma1, lambda1, gamma2, lambda2, U) {
  F_min_U <- 1 - exp(-(lambda1 * t^gamma1 + lambda2 * t^gamma2)) - U
  return(F_min_U)
}
# Generate uniform values and store in dataframe
u <- stats::runif(n)
# Inverse transform sampling
dat_roots <- cbind.data.frame(u, gamma1, lambda1, gamma2, lambda2)
times <- dat_roots %>%
  dplyr::rowwise() %>%
  dplyr::mutate(
    t_tilde = stats::uniroot(
      cdf_U,
      interval = c(.Machine$double.eps,
                    max_time),
      extendInt = "yes",
      U = u,
      gamma1 = gamma1,
      lambda1 = lambda1,
      gamma2 = gamma2,
      lambda2 = lambda2
    )$`root`
  ) %>%
  dplyr::pull(t_tilde)
# Generate event indicators
hazard1 <- weibull_hazard(gamma = gamma1, lambda = lambda1, t = times)
hazard2 <- weibull_hazard(gamma = gamma2, lambda = lambda2, t = times)
event <- stats::rbinom(n = n, size = 1, prob = hazard1 / (hazard1 + hazard2))
c.ind <- ifelse(event == 1, 1, 2)

```

```

# Add censoring
cens <- stats::rexp(n = n, rate = rate_cens)
c.ind <- ifelse(cens < times, 0, c.ind)
times <- pmin(cens, times)
# Winsorize time ranges to desired ones to make them more realistic
# and add some white noise
c.ind <- ifelse(times >= max_time, 0, c.ind)
times <- ifelse(times >= max_time, max_time, times)
times[times == max_time] <- times[times == max_time] +
  rnorm(length(times[times == max_time]), mean = 0, sd = 1e-4)
times <- ifelse(times < min_time, min_time, times)
times[times == min_time] <- times[times == min_time] +
  abs(rnorm(length(times[times == min_time]), mean = 0, sd = 1e-4))
sim.data <- data.frame(fstatus = c.ind, ftime = times)
X <- as.data.frame(X)
colnames(X) <- paste0("X", seq_len(p))
sim.data <- as.data.frame(cbind(sim.data, X))
return(sim.data)
}

gen_data <- function(n, p) {
  num_true <- 20
  beta1 <- c(rep(0, p))
  beta2 <- c(rep(0, p))
  nu_ind <- seq(num_true)
  # Here out of 20 predictors, 10 should be non-zero
  beta1[nu_ind] <- c(rep(1, 10), rep(0, 10))
  beta2[nu_ind] <- c(rep(-1, 10), rep(0, 10))

  # Simulate data
  sim.data <- cause_hazards_sim(n = n, p = p,
                                beta1 = beta1, beta2 = beta2,
                                rate_cens = 0.25,
                                h1 = 0.55, h2 = 0.05,
                                gamma1 = 1.5, gamma2 = 1.5,
                                exchangeable = TRUE)

  cen.prop <- c(prop.table(table(sim.data$fstatus)), 0, 0, 0, 0)

  # Training-test split
  # We only do this (instead of generating datasets for train and test
  # like Anthony mentioned because it is faster computationally
  # as casebase resamples) + proportion of censoring can be quite random
  # in each run of the simulation so we want to maintain the same in
  # validation and test set

  train.index <- caret::createDataPartition(sim.data$fstatus, p = 0.75, list = FALSE)

```

```

train <- sim.data[train.index,]
test <- sim.data[-train.index,]

#####
# We have two competitor models for variable selection:
# 1) Independent cox-regression model
# 2) penCR cox regression model - where the lambda penalties are trained together
##### Fit independent cox-regression model #####
##### Cause-1 #####
# Censor competing event
y_train <- Surv(time = train$ftime, event = train$fstatus == 1)

x_train <- model.matrix(~ . -ftime -fstatus, data = train)[, -1]

# Censor competing event
y_test <- Surv(time = test$ftime, event = test$fstatus == 1)

x_test <- model.matrix(~ . -ftime -fstatus, data = test)[, -1]

# Test set
surv_obj_val <- with(test, Surv(ftime, as.numeric(fstatus), type = "mstate"))

# Covariance matrix
cov_val <- cbind(test[, c(grepl("X", colnames(test)))], time = log(test$ftime))

# Case-base dataset
cb_data_val <- create_cbDataset(surv_obj_val, as.matrix(cov_val), ratio = 10)

return(list(X_train = cb_data_val$covariates,
            Y_train = cb_data_val$event_ind,
            offset = cb_data_val$offset,
            true_beta = cbind(beta1, beta2)))
}

```

2 Performance Evaluation —

In order to evaluate the performance of `MNlogistic`, we assess its convergence time and memory allocated with sample sizes of $n = 1000, 2000$ and covariates $p = 20, 100, 250, 500, 1000, 2000, 4000$. The original implementation is compared with an improved implementation `MNlogistic2` with the following changes:

- Matrix input `U` in proximal functions (Flat, Graph, Tree) are directly called using `U.memptr()`
- String `regul` in proximal functions (Flat, Graph, Tree) are directly called using `std::vector<char>`.
- Multinomial gradient functions changed to in-place calculation (takes `arma::mat&` `grad_out`, returns void), removing internal allocation.

- Multinomial gradient functions optimized to create column vector from x only once per call.

```

if (save){
  results <- bench::press(
    n = c(1000, 2000),
    p = c(20, 100, 250, 500, 1000, 2000, 4000),
    lambda = c(0.9, 0.5, 0.07, 0.005),
    {
      data <- gen_data(n, p)
      # lambda <- 0.9
      alpha = 0.5
      unpen_cov = 2
      # Elastic-net reparametrization
      lambda1 <- lambda*alpha
      lambda2 <- 0.5*lambda*(1 - alpha)
      # Prepare covariate matrix with intercept
      set.seed(1234)
      bench::mark(
        MNlogistic = mtool::mtool.MNlogistic(
          X = data$X_train,
          Y = data$Y_train,
          offset = data$offset,
          N_covariates = 2,
          regularization = 'elastic-net',
          transpose = FALSE,
          lambda1 = lambda1, lambda2 = lambda2,
          lambda3 = 0
        ),
        MNlogistic2 = mtool::mtool.MNlogistic2(
          X = data$X_train,
          Y = data$Y_train,
          offset = data$offset,
          N_covariates = 2,
          regularization = 'elastic-net',
          transpose = FALSE,
          lambda1 = lambda1, lambda2 = lambda2,
          lambda3 = 0
        ),
        glmnet = glmnet(data$X_train,
          data$Y_train,
          alpha = alpha,
          lambda = lambda,
          family = "multinomial"),
        min_iterations = 1, check = FALSE
      )
    }
  )
}

```

```

pushoverr::pushover(message = "Loop finished",
                    user = "uk2t3aqv1tavuxiuuzag7pf7742q15",
                    app = "aq5nckww93an2onvb4k1jin9487gak")

saveRDS(results, here::here("notes_jmr", "data", "results.rds"))
}

```

2.1 Plot results

2.1.1 Time

```

results <- readRDS(here::here("notes_jmr", "data", "results.rds"))

results %>%
  as.tibble() %>%
  unnest_longer(time) %>%
  ggplot(aes(x = p, y = time, group = paste0(expression, p, n),
            colour = as.character(expression))) +
  geom_jitter(alpha = 0.1) +
  stat_summary(geom = "point", fun.y = "mean", alpha = 1) +
  stat_summary(fun.data = mean_sdl, geom = "errorbar", width = 0.1) +
  stat_summary(aes(group = paste0(expression, n)), fun = mean, geom = "line") +
  facet_wrap(~paste0("n = ", n)) +
  labs(x = "p",
       y = "Time",
       colour = "Implementation")

```

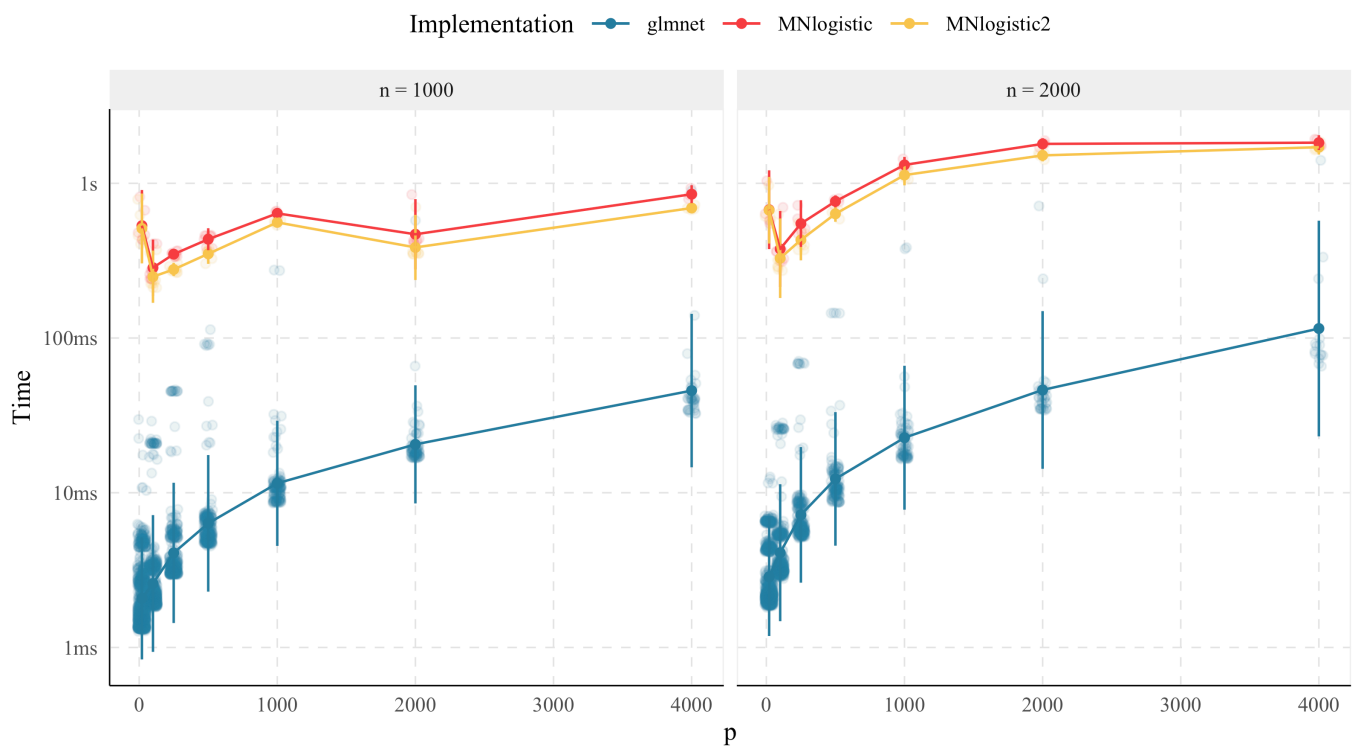
#> Warning: `as.tibble()` was deprecated in tibble 2.0.0.

#> i Please use `as_tibble()` instead.

#> i The signature and semantics have changed, see `?as_tibble`.

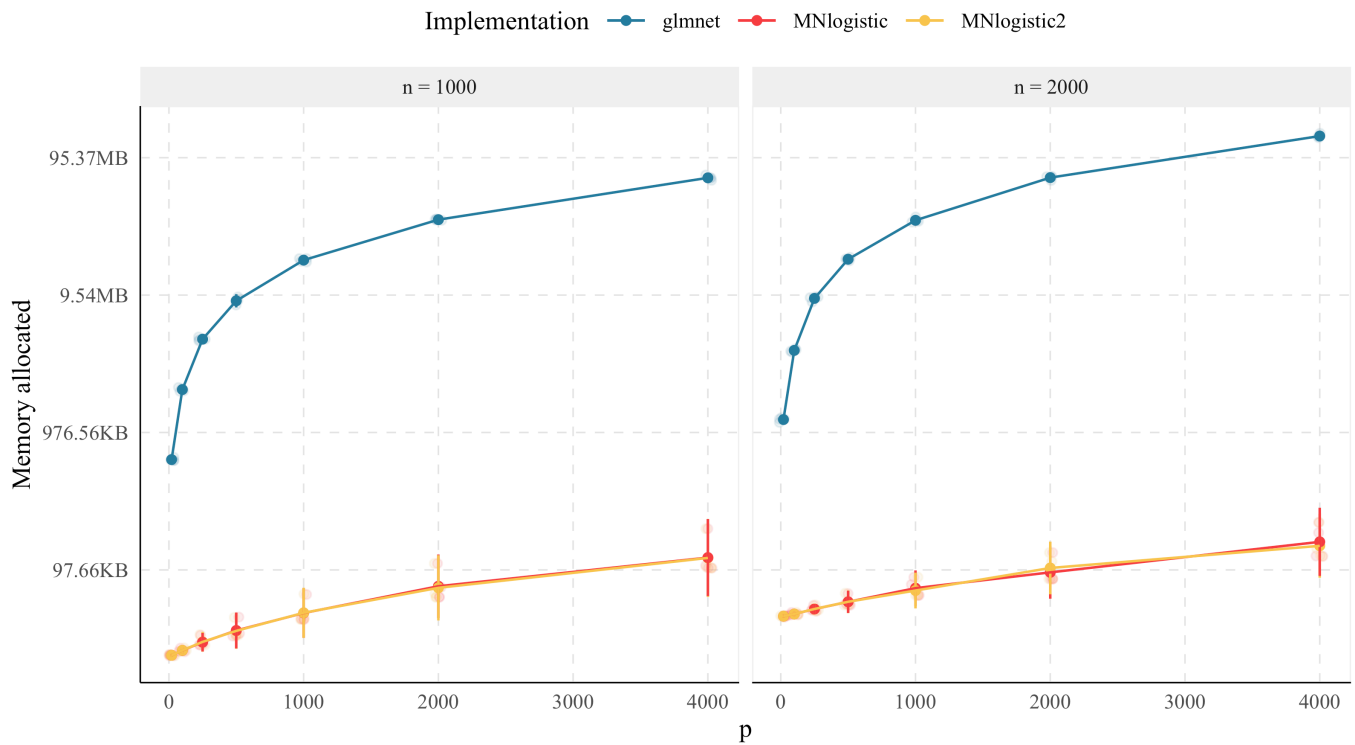
#> Warning: The `fun.y` argument of `stat_summary()` is deprecated as of ggplot2 3.3.0.

#> i Please use the `fun` argument instead.



2.1.2 Memory

```
results %>%
  as.tibble() %>%
  ggplot(aes(x = p, y = mem_alloc, group = paste0(expression, p, n),
             colour = as.character(expression))) +
  geom_jitter(alpha = 0.1) +
  stat_summary(geom = "point", fun.y = "mean", alpha = 1) +
  stat_summary(fun.data = mean_sdl, geom = "errorbar", width = 0.1) +
  stat_summary(aes(group = paste0(expression, n)), fun = mean, geom = "line") +
  facet_wrap(~paste0("n = ", n)) +
  labs(x = "p",
       y = "Memory allocated",
       colour = "Implementation")
```

2.2 Optimization Example

Objective function

```
objective <- function(B, X, Y, offset,
                      lambda, alpha, reg_p) {
```

```
  N_prime <- nrow(X)
```

```
  p <- ncol(X)
```

```
  K <- ncol(B)
```

```
  Y_int <- as.integer(Y)
```

```
  Y <- Y_int
```

```
  valid_y_indices <- which(Y >= 1 & Y <= K)
```

```
  if (length(valid_y_indices) < N_prime) {
```

```
    if (length(valid_y_indices) == 0) {
```

```
      loss_value <- 0.0
```

```
    penalty_value <- 0.0
```

```
    if (reg_p > 0 && lambda > 0) {
```

```
      B_reg <- B[1:reg_p, , drop = FALSE]
```

```
      lambda1 <- lambda * alpha
```

```
      lambda2 <- 0.5 * lambda * (1 - alpha)
```

```
      penalty_value <- (lambda1 * sum(abs(B_reg))) + (lambda2 * sum(B_reg^2))
```

```
    }
```

```
    return(penalty_value)
```

```
  }
```

```

    X_filt <- X[valid_y_indices, , drop = FALSE]
    Y_filt <- Y[valid_y_indices]
    offset_filt <- offset[valid_y_indices]
    N_filt <- length(valid_y_indices)
  } else {
    # Use all data if all Y are valid
    X_filt <- X
    Y_filt <- Y
    offset_filt <- offset
    N_filt <- N_prime
  }

  # eta = X %*% B + offset (N_filt x K)
  eta <- sweep(X_filt %*% B, 1, offset_filt, "+")

  # log(1 + sum_k(exp(eta_ik))) for each row i
  max_eta_stable <- apply(cbind(0, eta), 1, max) # max of 0 and row max(eta_k)
  log_one_plus_sum_exp <- max_eta_stable + log( exp(-max_eta_stable) + rowSums(exp(sweep(eta, 1, max_eta_stable, "-"))))

  idx_mat <- cbind(seq_len(N_filt), Y_filt)
  eta_yi <- eta[idx_mat]

  # -eta_yi + log(1 + sum_k exp(eta_ik))
  loss_per_obs <- -eta_yi + log_one_plus_sum_exp

  loss_value <- sum(loss_per_obs)

  # Penalty
  penalty_value <- 0.0
  if (reg_p > 0 && lambda > 0) {

    B_reg <- B[1:reg_p, , drop = FALSE]

    lambda1 <- lambda * alpha
    lambda2 <- 0.5 * lambda * (1 - alpha)

    l1_norm <- sum(abs(B_reg))

    l2_norm_sq <- sum(B_reg^2)

    penalty_value <- (lambda1 * l1_norm) + (lambda2 * l2_norm_sq)
  }

  return(loss_value + penalty_value)
}

```

```

if (save){
  set.seed(123)

  iter_sim <- tibble()

  for (i in 1) {
    n <- 1000
    p <- 20
    lambda = c(1e-10)
    data <- gen_data(n, p)
    alpha = 0.5
    unpen_cov = 2
    # Elastic-net reparametrization
    lambda1 <- lambda*alpha
    lambda2 <- 0.5*lambda*(1 - alpha)

    MNlogistic3 <- mtool::mtool.MNlogistic3(
      X = data$X_train,
      Y = data$Y_train,
      offset = data$offset,
      N_covariates = unpen_cov,
      regularization = 'elastic-net',
      transpose = FALSE,
      lambda1 = lambda1, lambda2 = lambda2,
      lambda3 = 0,
      niter_inner_mtplyr = 2, maxit = 100,
      tolerance = 1e-10
    )

    MNlogistic3$coefficients

    iter_vals <- map_dbl(MNlogistic3$coefficients$hist, ~
      objective(matrix(.x,
        nrow = ncol(data$X_train)),
        data$X_train, data$Y_train, data$offset,
        lambda, alpha, 20))

    iter_sim <- tibble(obj_step = iter_vals,
      iter = 1:length(iter_vals),
      sim = i) %>%
    bind_rows(iter_sim)
  }

  saveRDS(iter_sim, here::here("notes_jmr", "data", "opt-example.rds"))

}


iter_sim <- readRDS(here::here("notes_jmr", "data", "opt-example.rds"))

```

```

iter_sim %>%
  ggplot(aes(iter, y = obj_step, group = sim)) +
  geom_line(alpha = 0.3) +
  # geom_smooth(aes(group = 1)) +
  labs(y = "Obj. funct",
       x = "Iteration")

```



2.2.1 Comparison

```

if (save){
  for (i in 1:30) {

    seed <- as.integer(i)

    # take the last five digits of the initial seed
    the_seed = seed %% 100000

    set.seed(the_seed)
    n <- 1000
    p <- 2000
    lambda = c(1e-10)
    data <- gen_data(n, p)
    alpha = 0.5
    unpen_cov = 2
    # Elastic-net reparametrization
    lambda1 <- lambda*alpha
  }
}

```

```
lambda2 <- 0.5*lambda*(1 - alpha)
```

```
set.seed(the_seed)
time <- system.time({MNlogistic <- mtool::mtool.MNlogistic3(
  X = data$X_train,
  Y = data$Y_train,
  offset = data$offset,
  N_covariates = unpen_cov,
  regularization = 'elastic-net',
  transpose = FALSE,
  lambda1 = lambda1, lambda2 = lambda2,
  lambda3 = 0,
  niter_inner_mtplyr = 7, maxit = 100,
  tolerance = 1e-4
)})
```

```
time_tol <- system.time({MNlogistic_tol <- mtool::mtool.MNlogistic3(
  X = data$X_train,
  Y = data$Y_train,
  offset = data$offset,
  N_covariates = unpen_cov,
  regularization = 'elastic-net',
  transpose = FALSE,
  lambda1 = lambda1, lambda2 = lambda2,
  lambda3 = 0,
  niter_inner_mtplyr = 7, maxit = 100,
  tolerance = 1e-5
)})
```

```
set.seed(the_seed)
timeExp <- system.time({MNlogisticExp <- mtool::mtool.MNlogisticExp(
  X = data$X_train,
  Y = data$Y_train,
  offset = data$offset,
  N_covariates = unpen_cov,
  regularization = 'elastic-net',
  transpose = FALSE,
  lambda1 = lambda1, lambda2 = lambda2,
  lambda3 = 0,
  niter_inner_mtplyr = 2, maxit = 100,
  tolerance = 1e-4,
  learning_rate = 1e-3
)})
```

```
set.seed(the_seed)
timeAcc <- system.time({MNlogisticAcc <- mtool::mtool.MNlogisticAcc(
```

```

X = data$X_train,
Y = data$Y_train,
offset = data$offset,
N_covariates = unpen_cov,
regularization = 'elastic-net',
transpose = FALSE,
lambda1 = lambda1, lambda2 = lambda2,
lambda3 = 0,
niter_inner_mtplyr = 1,
maxit = 50,
momentum_gamma = .9,
tolerance = 1e-1,
learning_rate = 1e-4,
pos = T
)))

```

```

set.seed(the_seed)
timeSAHRA <- system.time({MNlogisticSAHRA <- mtool::mtool.MNlogisticSAHRA(
  X = data$X_train,
  Y = data$Y_train,
  offset = data$offset,
  N_covariates = unpen_cov,
  regularization = 'elastic-net',
  transpose = F,
  lambda1 = lambda1, lambda2 = lambda2,
  lambda3 = 0,
  niter_inner_mtplyr = 1,
  maxit = 50,
  tolerance = 1e-4,
  learning_rate = 1e-3
)})

```

```

iter_vals <- map_dbl(MNlogistic$coefficientshist, ~
  objective(matrix(.x,
    nrow = ncol(data$X_train)),
    data$X_train, data$Y_train, data$offset,
    lambda, alpha, 20))

iter_vals_tol <- map_dbl(MNlogistic_tol$coefficientshist, ~
  objective(matrix(.x,
    nrow = ncol(data$X_train)),
    data$X_train, data$Y_train, data$offset,
    lambda, alpha, 20))

iter_valsExp <- map_dbl(MNlogisticExp$coefficientshist, ~
  objective(matrix(.x,
    nrow = ncol(data$X_train)),
    data$X_train, data$Y_train, data$offset,

```

```

        lambda, alpha, 20))

iter_valsAcc <- map_dbl(MNlogisticAcc$coefficientshist, ~
  objective(matrix(.x,
    nrow = ncol(data$X_train)),
    data$X_train, data$Y_train, data$offset,
    lambda, alpha, 20))

iter_valsSAHRA <- map_dbl(MNlogisticSAHRA$coefficientshist, ~
  objective(matrix(.x,
    nrow = ncol(data$X_train)),
    data$X_train, data$Y_train, data$offset,
    lambda, alpha, 20))

# MNlogistic$coefficients
# MNlogisticExp$coefficients
  if (i == 1) iter_sim <- tibble()

iter_sim <- tibble(obj_step = iter_vals,
  iter = 1:length(iter_vals),
  sim = i,
  time = time[1],
  algorithm = "Original implementation\n(tol = 1e-4)") %>%
bind_rows(tibble(obj_step = iter_vals_tol,
  iter = 1:length(iter_vals_tol),
  sim = i,
  time = time_tol[1],
  algorithm = "Original implementation \n(tol = 1e-5)"),
  tibble(obj_step = iter_valsExp,
  iter = 1:length(iter_valsExp),
  sim = i,
  time = timeExp[1],
  algorithm = "MNlogisticExp \n(Adj. param)"),
  tibble(obj_step = iter_valsAcc,
  iter = 1:length(iter_valsAcc),
  sim = i,
  time = timeAcc[1],
  algorithm = "Acc. MNlogisticExp"),
  tibble(obj_step = iter_valsSAHRA,
  iter = 1:length(iter_valsSAHRA),
  sim = i,
  time = timeSAHRA[1],
  algorithm = "SAHRA")) %>%
  bind_rows(iter_sim)
saveRDS(iter_sim, here::here( "notes_jmr", "data", "opt-algorithms.rds"))
}

```

```

}

iter_sim <- readRDS(here::here("notes_jmr", "data", "opt-algorithms.rds"))

iter_sim %>%
  ggplot(aes(iter, y = obj_step, group = paste0(algorithm, sim),
            color = str_wrap(algorithm, 15))) +
  # geom_line(alpha = 0.2) +
  stat_summary(aes(group = algorithm), geom = "point", fun.y = "mean", alpha = 0.5) +
  stat_summary(aes(group = algorithm), fun = mean, geom = "line", alpha = 0.5) +
  labs(y = "Obj. funct",
       x = "Epoch",
       color = "") +
  xlim(0,9)

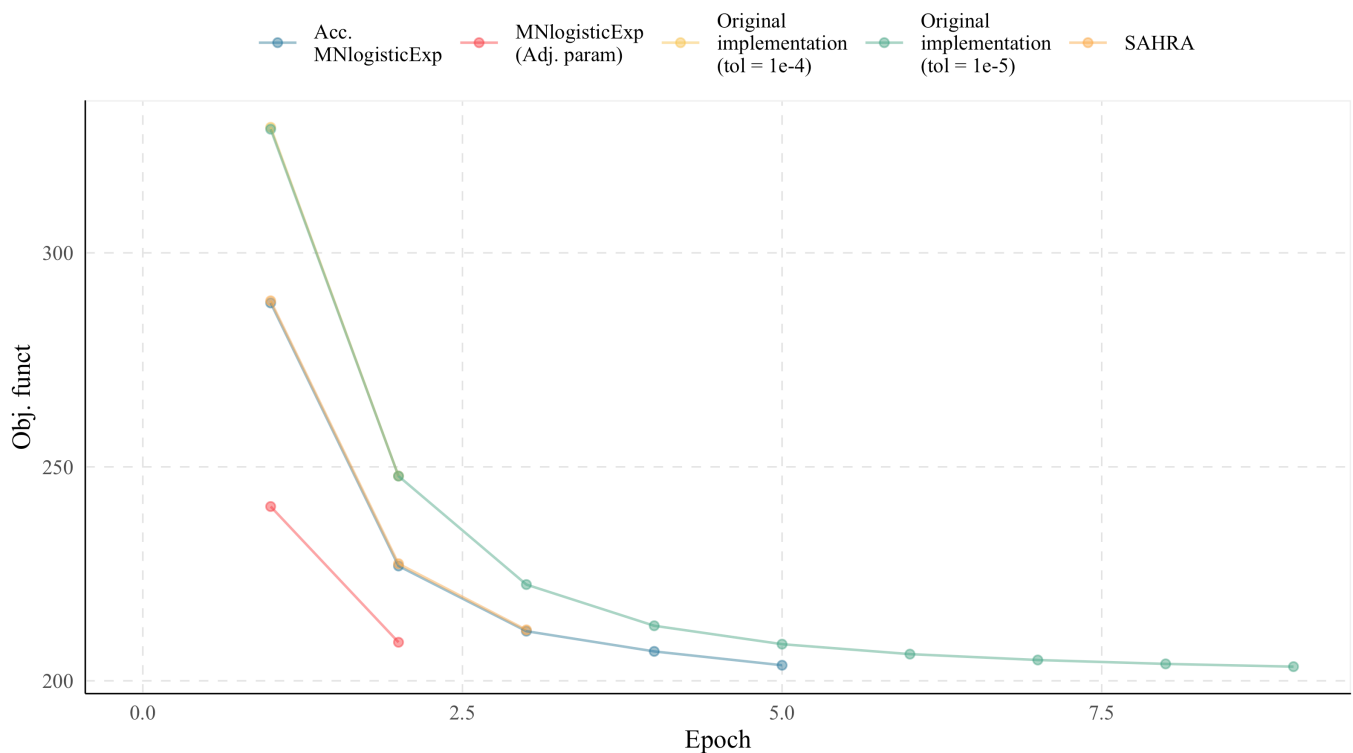
```

#> Warning: Removed 4 rows containing non-finite outside the scale range

#> (`stat_summary()`).

#> Removed 4 rows containing non-finite outside the scale range


#> (`stat_summary()`).



```

iter_sim %>%
  group_by(sim, algorithm) %>%
  summarise(time = first(time)) %>%
  ggplot(aes(x = algorithm, y = time)) +
  geom_boxplot() +
  labs(y = "Time",
       x = "Algorithm")

```

Notes:

- Inner loop iterations do not appear to increase the number of epochs proportionally, suggesting that it may reduce computational time. The inner loop iterations set to $2n$ appear to perform effectively, reducing computational time to $1/3$.
 - As the value of p increases, it seems that the algorithm fails to reach the minimum.
 - The situation becomes problematic when $p \gg n$ because the initial epoch is below the tolerance level and far from the optimal position. To address this issue, we can lower the tolerance; however, this increases the number of steps required. By reducing the tolerance, we can observe that it becomes necessary to achieve reasonable coefficients in this context.
- **Tolerance** of $1e-4$ in high-dimensional settings is not enough to reach optimal levels.
 - Maybe, we can reduce the number of iterations as we increase the number of epochs.
 - Under similar tolerance, a smaller tolerance seems to be better.
- **Learning rate.** The increase in learning rate from $1e-4$ to $1e-3$ reduces computational time by half. We should consider better learning rates. A smaller learning rate generates NaNs in the Frobenius norm.